black hat
USA 2022

AUGUST 10-11, 2022

BRIEFINGS

Trace Me If You can:
Bypassing Linux Syscall Tracing

Dr. Rex Guo, Lacework
Dr. Junyuan Zeng, Linkedin.com

About Rex Guo

- Principal Engineer @ Lacework
  - Behavior anomaly detection (Polygraph)
  - CSPM
- Engineering Manager @ Startups
  - Confluera (XDR)
  - Tetration (CWPP, now part of Cisco)
- Conference speaker at Blackhat, DEFCON,...
- @Xiaofei_REX

LACEWORK

About Junyuan Zeng

- Linkedin
  - Senior Software Engineer: Kubernetes

- JD.com
  - Staff Security Architect/Engineer: Cloud native security

- Samsung Research America & FireEye
  - Staff Security Software Engineer/Researcher: Mobile security

An Incident - An Attacker's View

Log4shell RCE on joe-box and executed a reverse shell

LACEWORK

An Incident - An Attacker's View

Log4shell RCE on joe-box and executed a reverse shell

Privilege escalation using sudo CVE-2021-3156

LACEWORK

An Incident - An Attacker's View

Log4shell RCE on
joe-box and
executed a reverse
shell

read /etc/shadow

Privilege escalation
using sudo
CVE-2021-3156

LACEWORK

An Incident - An Attacker's View

Log4shell RCE on
joe-box and
executed a reverse
shell

read /etc/shadow

Privilege escalation
using sudo
CVE-2021-3156

Read ssh process
environment
variable

LACEWORK

An Incident - An Attacker's View

Log4shell RCE on
joe-box and
executed a reverse
shell

read /etc/shadow

Lateral movement
to alice-box with
ssh hijacking

Privilege escalation
using sudo
CVE-2021-3156

Read ssh process
environment
variable

LACEWORK

An Incident - An Attacker's View

Log4shell RCE on joe-box and executed a reverse shell

read /etc/shadow

Lateral movement to alice-box with ssh hijacking

Privilege escalation using sudo CVE-2021-3156

Read ssh process environment variable

?

LACEWORK

An Incident - A Defender's View

Log4shell RCE on joe-box and executed a reverse shell

read /etc/shadow

Lateral movement to alice-box with ssh hijacking

Privilege escalation using sudo CVE-2021-3156

Read ssh process environment variable

connect etc.

execve(at)

open(at)

open(at)

connect, etc.

Incident Response
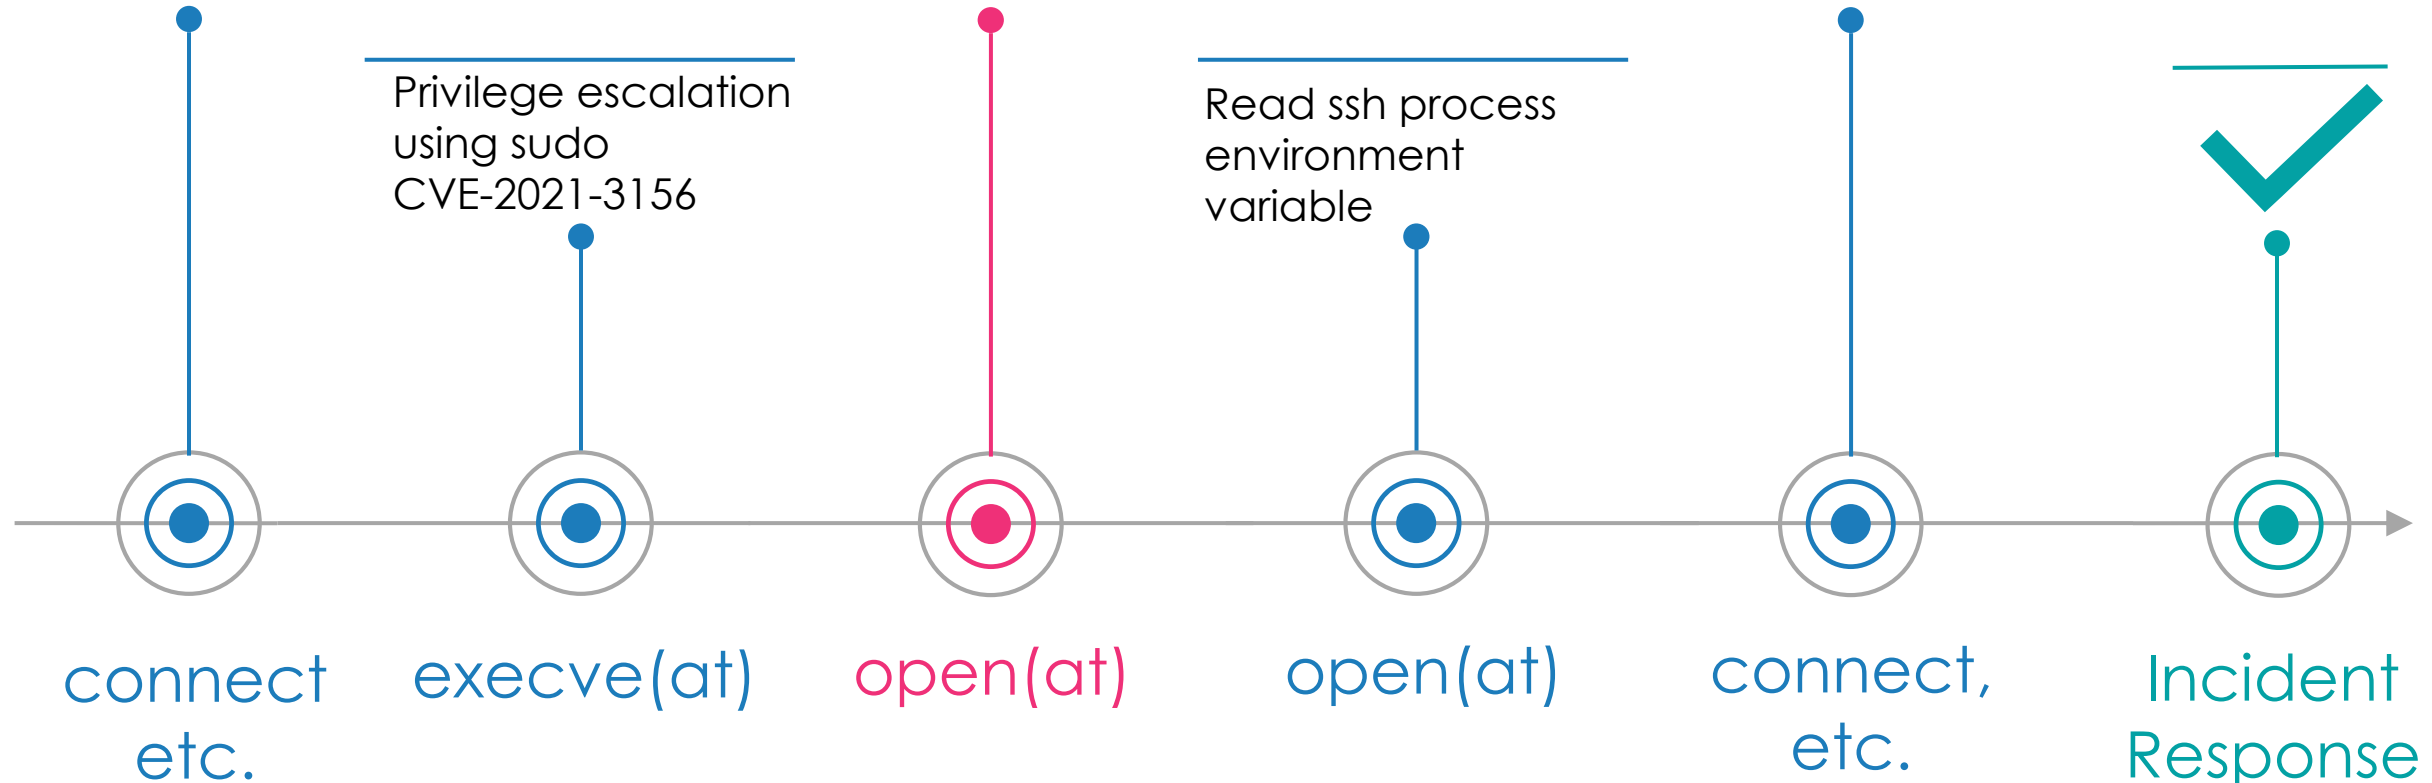
Detection Rule Example

rule: untrusted program reads */etc/shadow*

condition:

    syscall == open(at)

    and has read permission

    and filename == */etc/shadow*

    and program is not in allowlist

Agenda

- Syscall Tracing
- Vulnerabilities
- Exploitations
- Mitigations
- Takeaways

LACEWORK

System Call Tracing
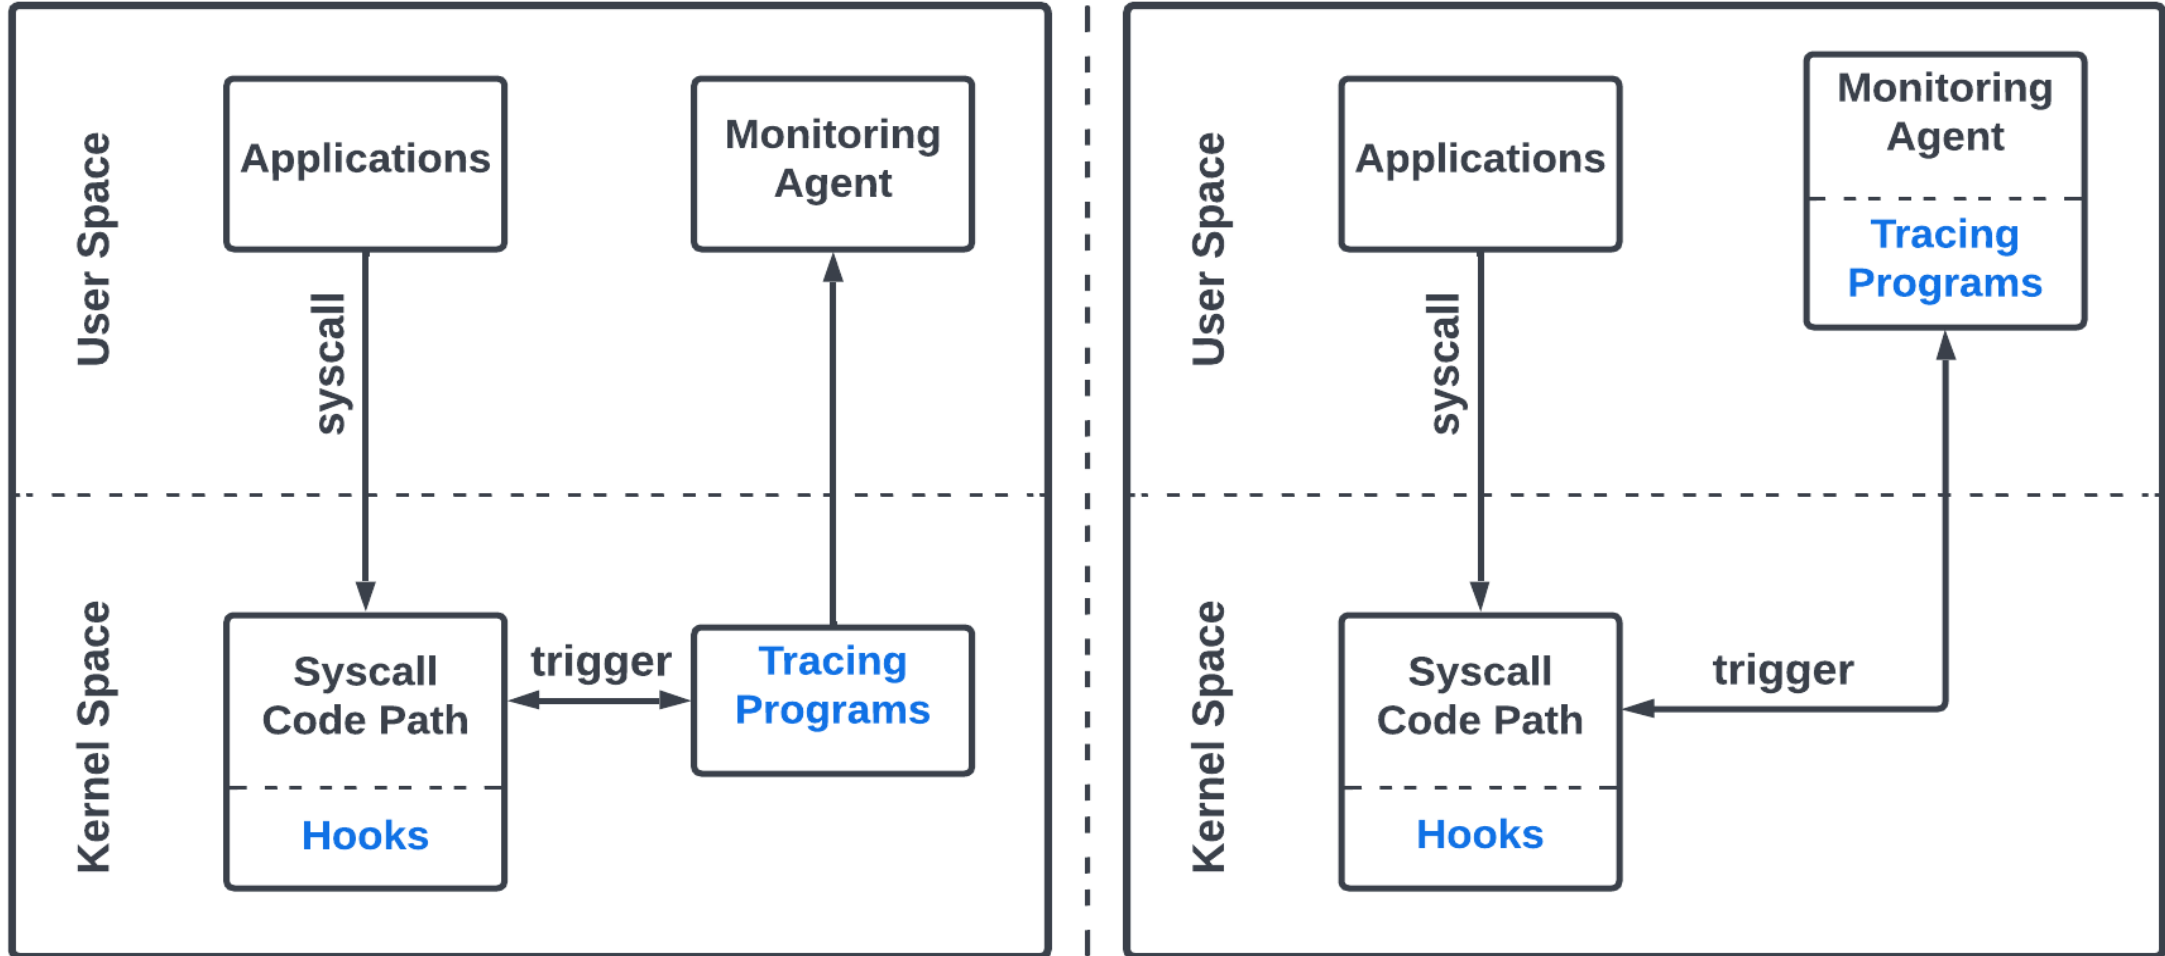
System Call Tracing – Tracing Program

- Tracing programs collect system call data, e.g., arguments

- Tracing programs can "attach" to different hooks
  - tracepoints, kprobe, ptrace etc.

- Tracing programs implementations
  - Linux native mechanisms: *ftrace*, *perf_events* etc.
  - Kernel modules, eBPF probe and user space programs

System Call Tracing – tracepoint

- **tracepoint**
  - Kernel static hook
  - Linux kernel provides *sys_enter* and *sys_exit*
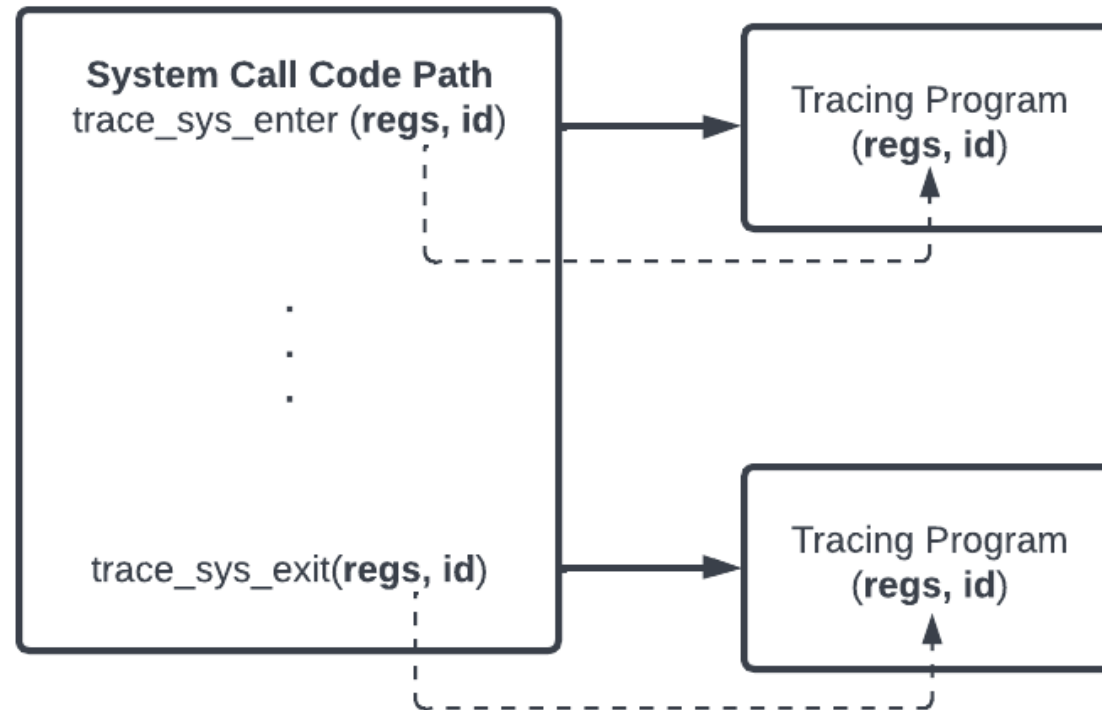
  - Low overhead but only static interceptions

System Call Tracing – kprobe

- kprobe
  - Dynamic hook in the kernel
  - Register tracing programs on instructions in syscall code path
  - Dynamic but slow compared to tracepoint and need to know exactly how data is placed on the stack and registers

System Call Code Path

Inst. 1
Inst. 2

.
.
.

Inst. n

Tracing Program
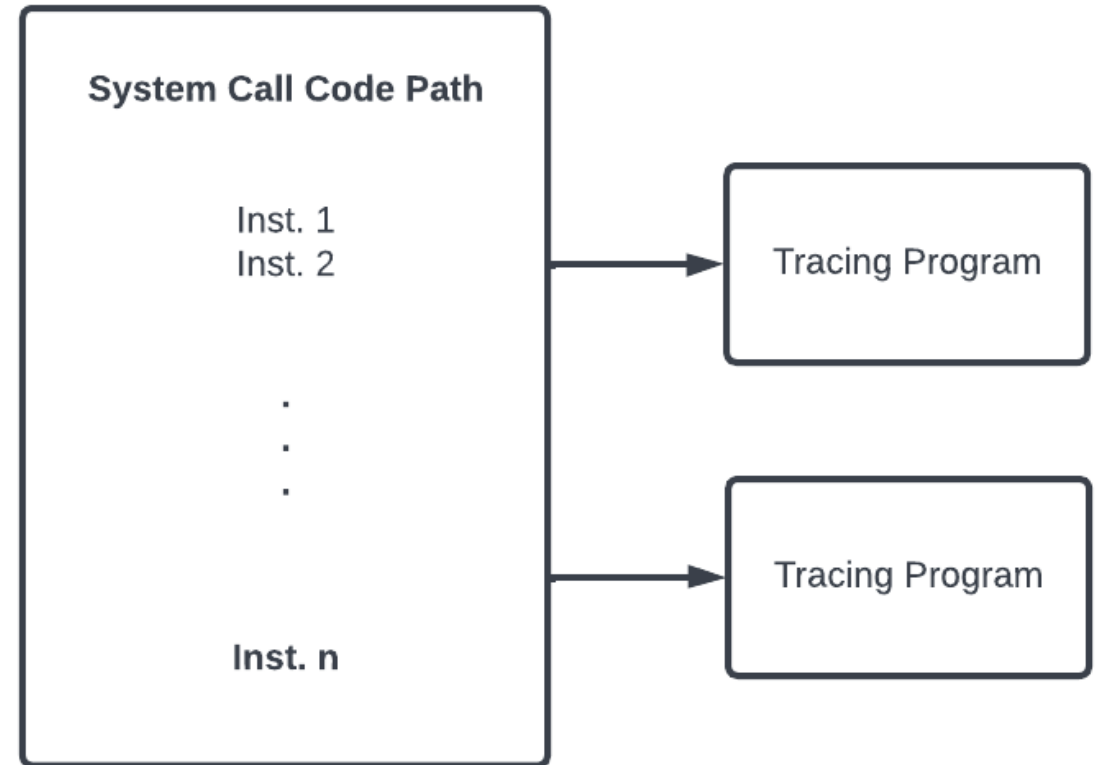
Tracing Program

System Call Tracing – ptrace

- ptrace
  - A static hook
  - No Kernel Module/eBPF program are needed
  - Performance overhead is high
  - Can combine with seccomp to reduce overhead
- Others (LD_PRELOAD etc.)

Cloud Workloads

- Virtual machines
  - AWS EC2 instances
  - Google VM instances
- Containers on customer-managed VMs
  - AWS EC2 tasks
  - Standard GKE workloads (e.g. DaemonSet etc.)
  - AKS workloads
- Serverless containers: have no access to the host
  - AWS Fargate tasks
  - GCP Cloud Run services
- Others (AWS Lambda etc.)

LACEWORK

System Call Tracing for Cloud Workloads

Workload	System Call Tracing
VMs	• Hooks: tracepoint, kprobe, ptrace • Tracing programs: kernel programs (eBPF, kernel Module), user programs • Tools: Falco eBPF/kernel Module, Falco pdig
Containers	• Hooks: tracepoint, kprobe, ptrace • Tracing programs: kernel programs (eBPF, kernel Module), user programs • Tools: Falco eBPF/kernel Module, Falco pdig
Serverless Containers	• Hooks: ptrace • Tracing programs: user programs • Tools: Falco pdig

LACEWORK

Open Source Projects

- Falco
  - Open source endpoint security monitoring project in CNCF
  - 5K+ github stars
  - Falco supports syscall tracing techniques:
    - tracepoint + kernel module
    - tracepoint + eBPF probe
    - pdig: ptrace + userspace program

- Falco pdig
  - Support syscall tracing of serverless workloads

LACEWORK

TOCTOU in Syscall Tracing

- *sys_connect(int fd, struct sockaddr __user * uservaddr, int addrlen)*

- TOC (Time-Of-Check): tracing programs dereference this user space pointer

- TOU (Time-Of-Use): the kernel dereferences this user space pointer

TOCTOU - Connect

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

ptrace/seccomp/sysenter tracepoint

User Space
Kernel 5.7.0

Execution Flow

```
Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…
```

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
    {
        …
        struct filename *tmp;
        ret = move_addr_to_kernel
                (uservaddr, addrlen, &address);
        if (!ret)
          ret = __sys_connect_file
                (f.file, &address, addrlen, 0);

        …
    }
```

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

TOCTOU - Connect

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

User Space
Kernel 5.7.0

Execution Flow

Syscall Table (x86_64)

…

42 sys_connect

43 sys_accept

44 sys_sendto

…

```
long __sys_connect((int fd,
      struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
      ret = __sys_connect_file
            (f.file, &address, addrlen, 0);
    …
}
```

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

LACEWORK

TOCTOU - Connect

syscall enter

User Space Kernel 5.7.0 | **Execution Flow**

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

Syscall Table (x86_64)

…

42 sys_connect

43 sys_accept

44 sys_sendto

…

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
      ret = __sys_connect_file
            (f.file, &address, addrlen, 0);

    …
}
```

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

ptrace/sysexit tracepoint

syscall exit

TOCTOU - Connect

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

■ **Userspace pointer pointing to "socket address"**

User Space
Kernel 5.7.0

Execution Flow

Syscall Table (x86_64)

…

42 sys_connect

43 sys_accept

44 sys_sendto

…

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
      ret = __sys_connect_file
            (f.file, &address, addrlen, 0);

    …
}
```

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

LACEWORK

TOCTOU - Connect

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

■ **Userspace pointer pointing to "socket address"**

■ **Kernel pointer pointing to "socket address"**

Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…

User Space
Kernel 5.7.0

Execution Flow

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
    {
        …
        struct filename *tmp;
        ret = move_addr_to_kernel
                (uservaddr, addrlen, &address);
        if (!ret)
          ret = __sys_connect_file
                (f.file, &address, addrlen, 0);

        …
    }
```

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

LACEWORK.

TOCTOU - Connect

User Space
Kernel 5.7.0

Execution Flow

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
      ret = __sys_connect_file
            (f.file, &address, addrlen, 0);

    …
}
```

**TOU by
Linux Kernel**

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

syscall exit

LACEWORK

TOCTOU - Connect

syscall enter →

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…

User Space
Kernel 5.7.0

Execution Flow

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
    {
        …
        struct filename *tmp;
        ret = move_addr_to_kernel
                (uservaddr, addrlen, &address);
        if (!ret)
          ret = __sys_connect_file
                (f.file, &address, addrlen, 0);

        …
    }
```

TOU by
Linux Kernel

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

syscall exit ←

LACEWORK

TOCTOU - Connect

syscall enter

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

sys_enter tracepoint
ptrace

TOC by Tracing Programs

kprobe

Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…

User Space
Kernel 5.7.0

Execution Flow

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
        ret = __sys_connect_file
            (f.file, &address, addrlen, 0);
    …
}
```

TOU by Linux Kernel

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

syscall exit

TOCTOU - Connect

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

Syscall Table (x86_64)

…

42 sys_connect
43 sys_accept
44 sys_sendto

…

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
        ret = __sys_connect_file
                (f.file, &address, addrlen, 0);

    …
}
```

**TOU by
Linux Kernel**

**User Space
Kernel 5.7.0**

Execution Flow

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

**sys_exit tracepoint
ptrace**

**TOC by Tracing
Programs**

TOCTOU - Connect

syscall enter →

User Space
Kernel 5.7.0

Execution Flow

syscall exit

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

Falco pdig

```
Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…
```

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
      ret = __sys_connect_file
            (f.file, &address, addrlen, 0);

    …
}
```

TOU by Linux Kernel

```
trace_sys_exit(regs, regs->ax)
```

Falco (<0.31.1) kernel module/eBPF

```
ptrace_report_syscall(regs, message)
```

Falco pdig

TOCTOU Windows across Kernels

- TOCTOU windows exist since the initial release of tracepoint/ptrace
- Expected behaviors
- Monitor kernel memory

TOCTOU – Falco

- User space pointers are dereferenced by
  - *sys_exit* tracepoint (kernel module, eBPF)
  - *sys_exit* ptrace (pdig)
- Falco older than v0.31.1
  - Check with vendors which commercial versions are affected
- 12/06/2021 Issue reported (CVE-2022-26316)
- 03/11/2022 Mitigation implemented (Advisory)
  - For selected syscalls, compare sys_enter and sys_exit tracepoint data (Falco LKM, eBPF)
  - Compare sys_enter and sys_exit ptrace data (Falco pdig)

LACEWORK

TOCTOU – Falco

- We evaluated the important syscalls in [Falco rules](#).

Syscall	Category	TOCTOU?	Exploitable by blocking condition	Exploitable by DC29 attack
connect	Network	Y	Y	Y
sendto/sendmsg	Network	Y	N	Y
open(at)	File	Y	Y	Y
execve	File	N	N*	N*
rename	File	Y	Y	Y
renameat(2)	File	Y	Y	Y
mkdir(at)	File	Y	Y	Y
rmdir	File	Y	Y	Y
unlink(at)	File	Y	Y	Y
symlink(at)	File	Y	Y	Y
chmod/fchmod(at)	File	Y	Y	Y
creat	File	Y	Y	Y

Exploit Requirements

| tracing program read register @ sys_enter TOC	How to increase delay? →	syscall code copy_from_user TOU	How to increase delay? →	tracing program read register @ sys_exit TOC

- Exploitation requirements
  - No additional privilege and capabilities
  - Control the time to inject the delay
 - Enough delay for pointer overwrite
  - Reliable

Exploit Strategy 1 (DEFCON 29)

- Injected delay is small

- Requires Userfaultfd syscall for precise injection while pausing the kernel execution

- seccomp can block userfaultfd syscall (e.g., docker default seccomp profile)

- Falco's mitigation was to detect userfaultfd
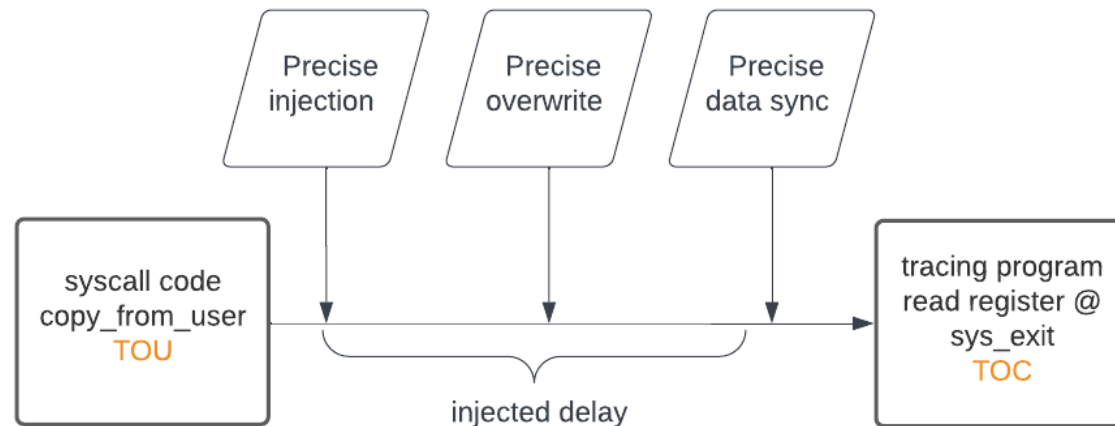
Exploit Strategy 2

- Injected delay >> the syscall execution time
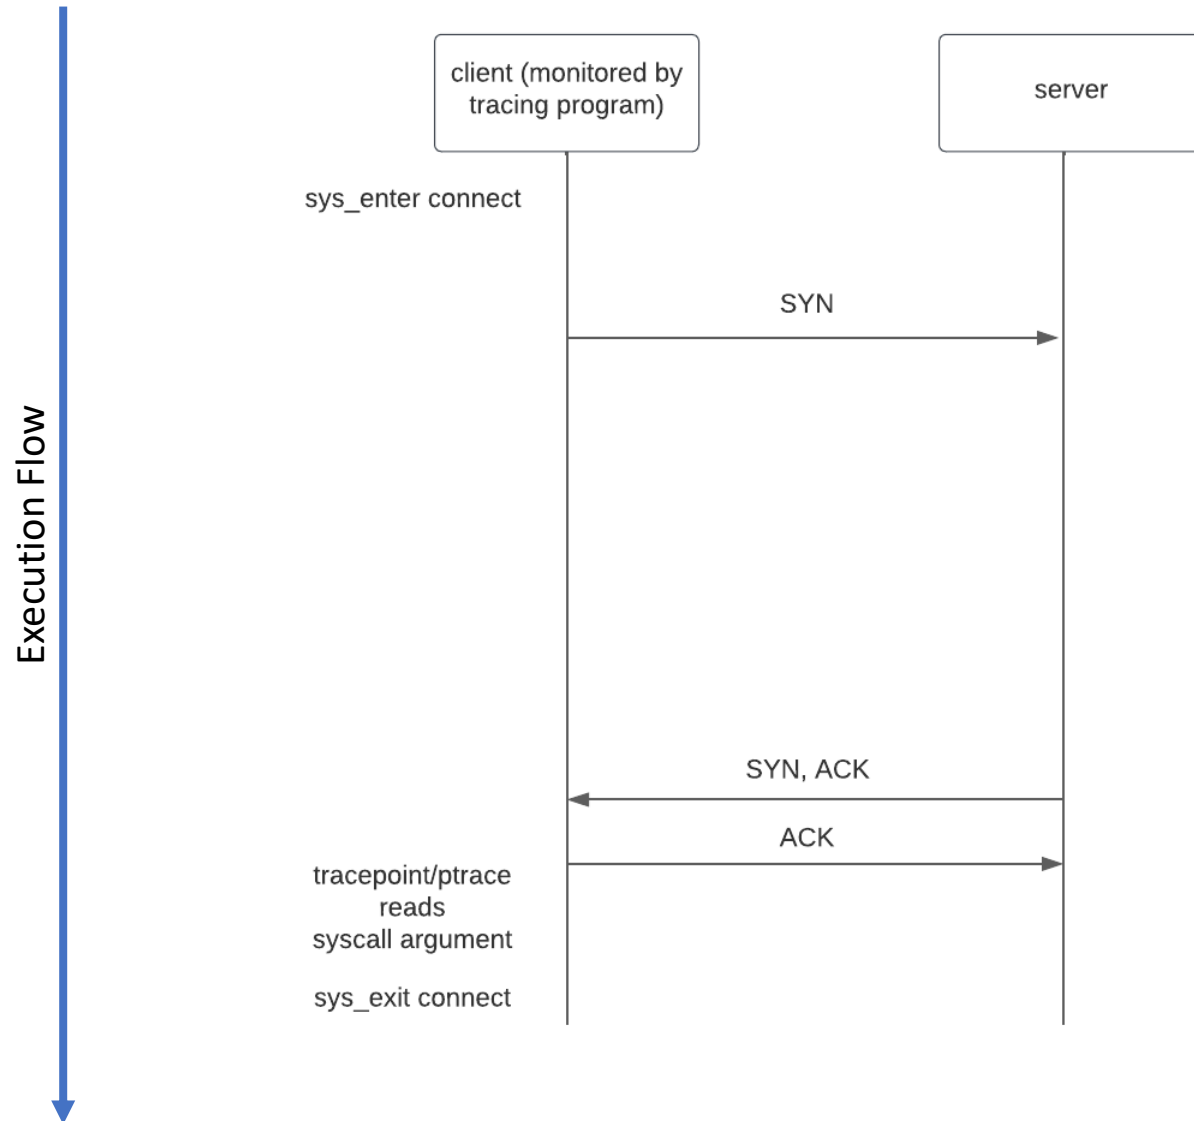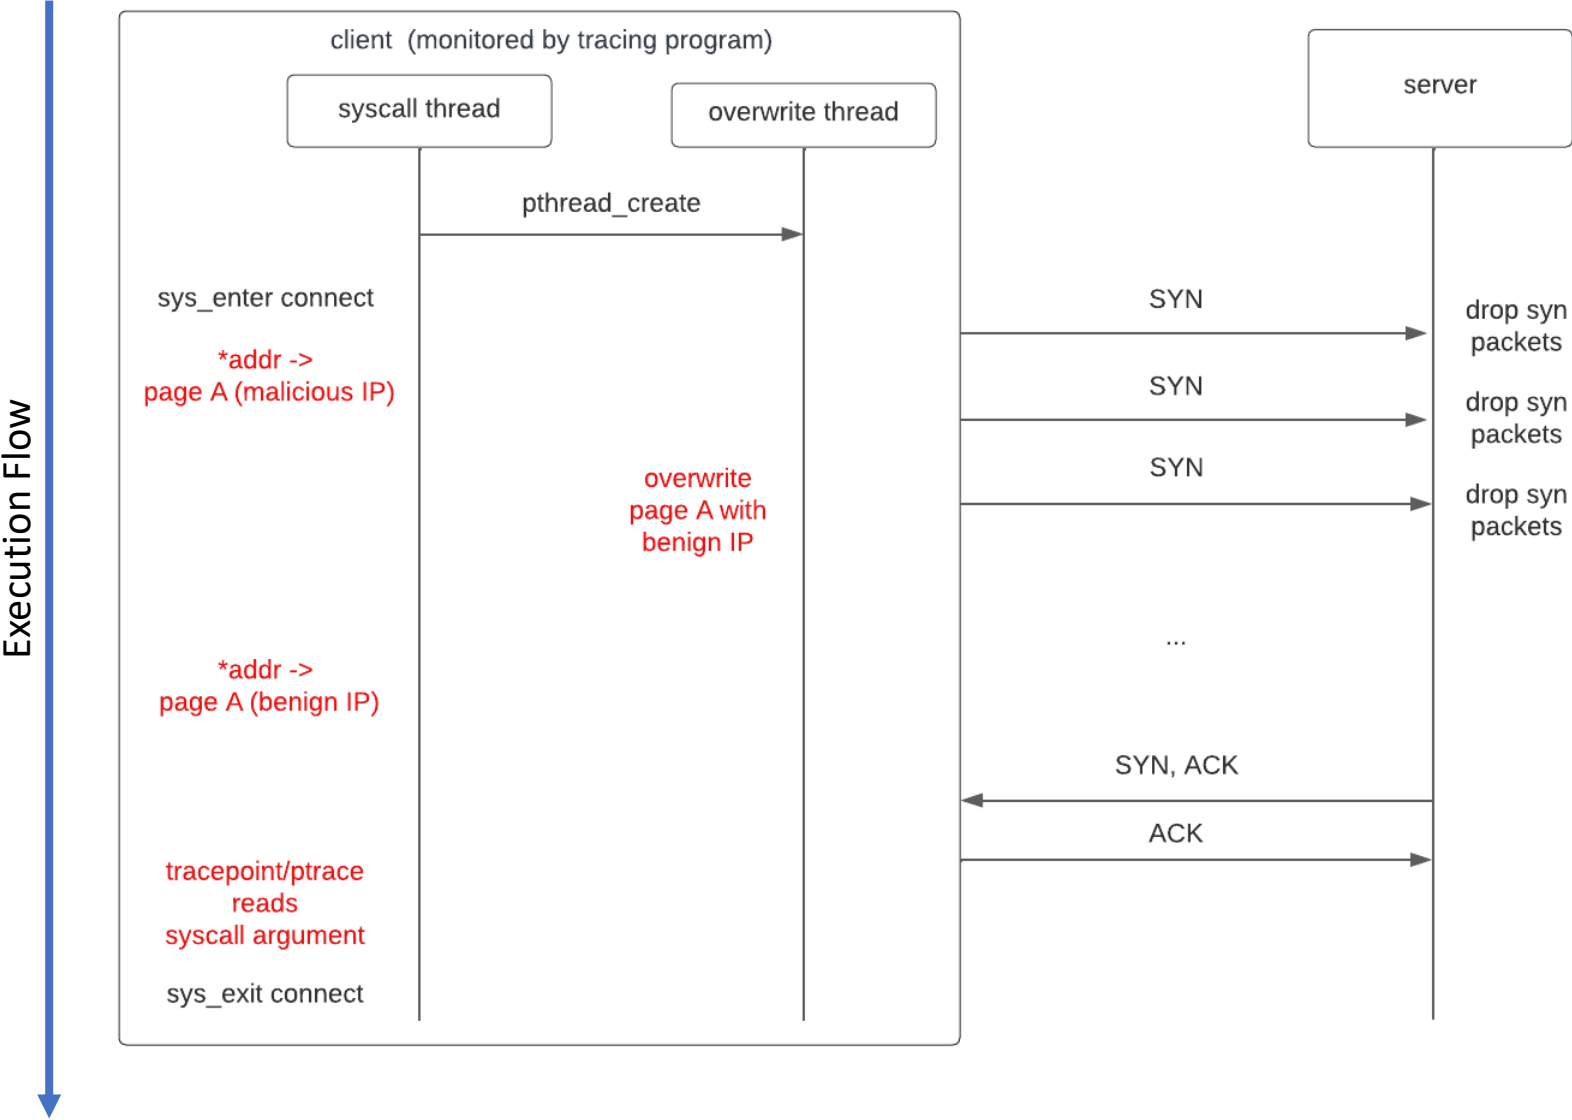- No Precise control is required

Syscall Built-in Delay

- Attackers can trigger significant syscall delays by introducing:
  - Blocking conditions (attack sys_exit)
  - Seccomp rules (attack sys_enter)
- Syscall can get "blocked"

Categories	Syscalls
Process	fork/exec/exit/wait/…
File system	open(at)/symlink(at)/read/write/…
Networking	connect/accept/socket/…
Security	seccomp/keyctl/…
Many others…	…

Connect Syscall

Execution Flow

client (monitored by tracing program)   server

sys_enter connect

SYN

SYN, ACK

ACK

tracepoint/ptrace
reads
syscall argument

sys_exit connect

LACEWORK

Bypassing Connect Syscall Tracing (Demo)

Blocking Syscalls (File Systems)

- File system syscalls are all affected
  - open/openat
  - creat
  - rename/renameat/renameat2
  - mkdir/mkdirat
  - rmdir
  - Other file system syscalls with pointer arguments

- Other syscalls are also affected due to fetching files from file systems.
  - execve/execveat

LACEWORK.

Filesystem in USErspace - FUSE

- User space filesystem framework

- Used as remote storage FUSE
  - Access the remote files as local ones
  - Faster evolvement and don't panic the kernel etc.

- Remote storage FUSE examples:
  - gcsfuse[1] : developed by Google for GCS
  - s3fs-fuse[2]: Amazon S3
  - BlobFuse[3]: developed by Azure for Blob storage
  - MezzFS[4]: developed and deployed @ Netflix
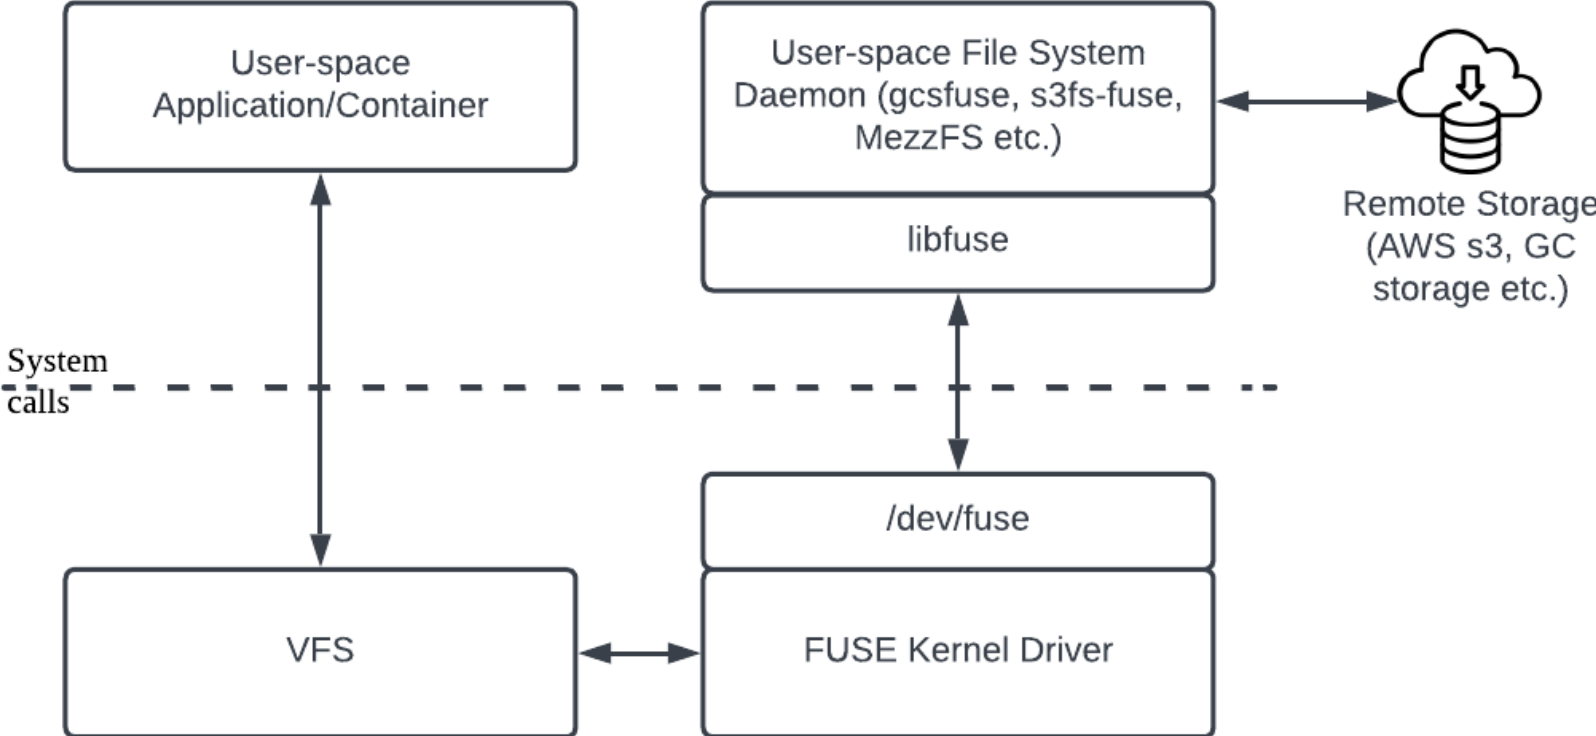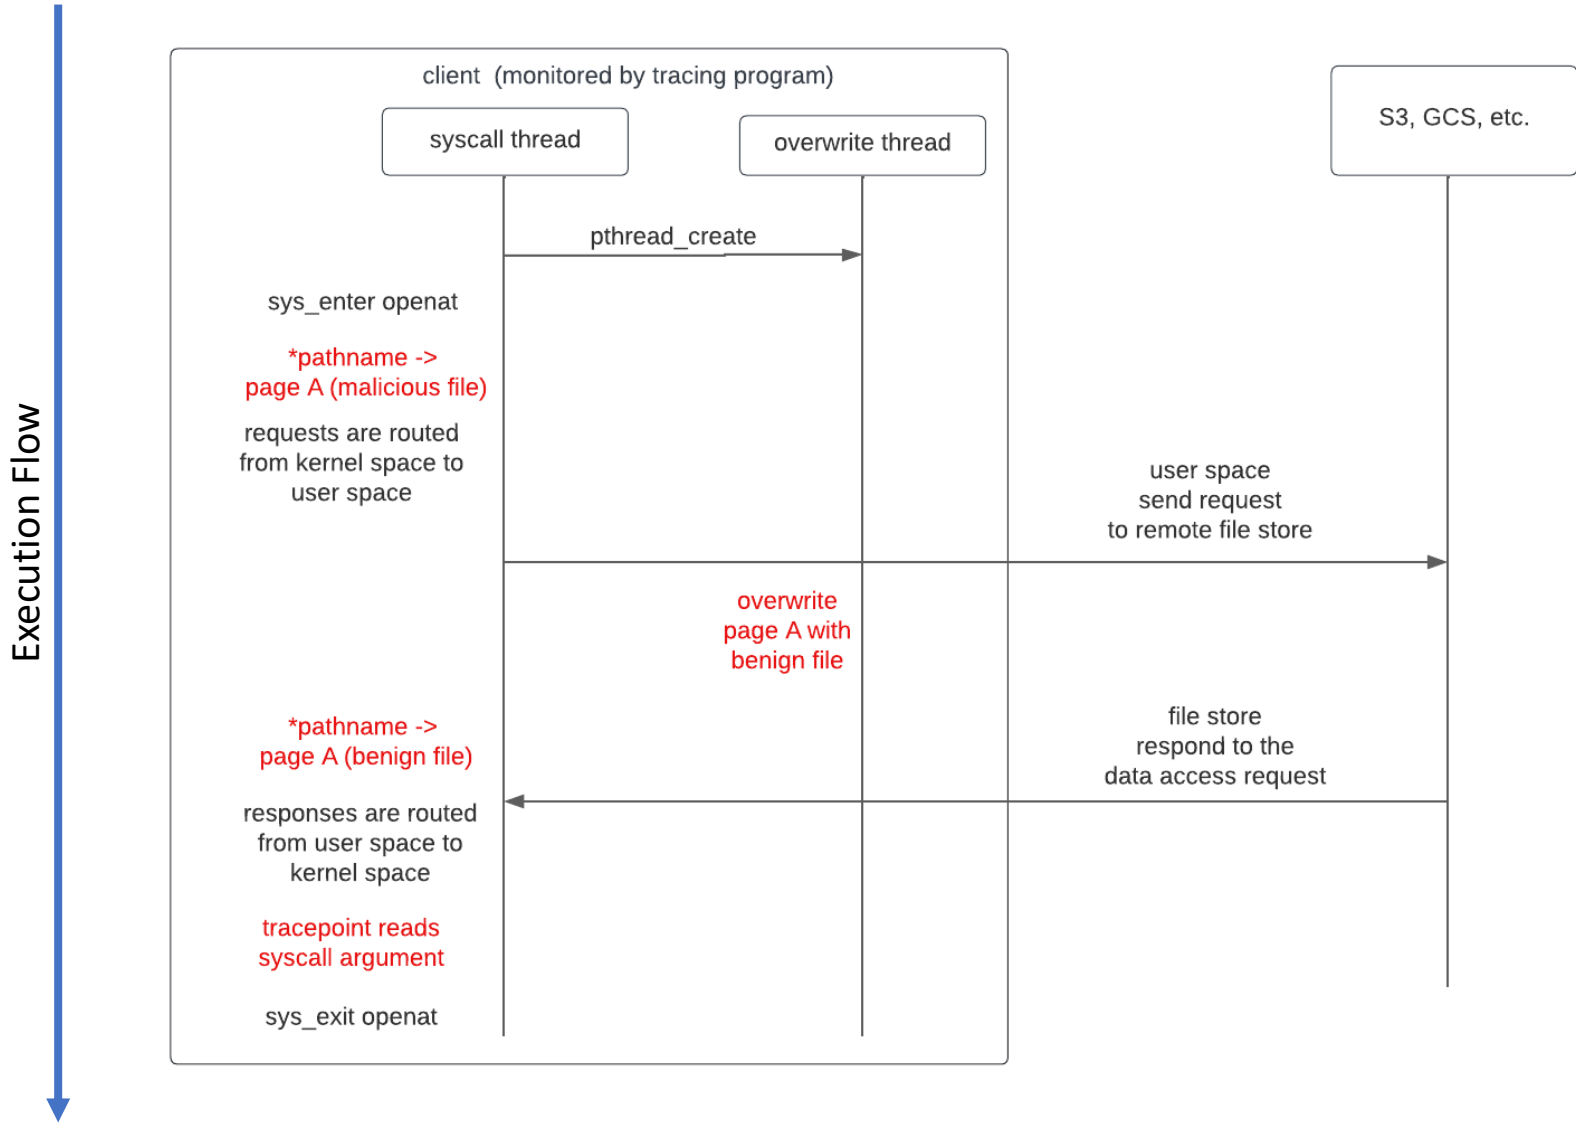 - Many others (sshfs etc.)

1, https://github.com/GoogleCloudPlatform/gcsfuse
2, https://github.com/s3fs-fuse/s3fs-fuse
3, https://github.com/Azure/azure-storage-fuse
4, https://netflixtechblog.com/mezzfs-mounting-object-storage-in-netflixs-media-processing-platform-cda01c446ba

LACEWORK

Remote Storage FUSE - Architecture

Bypassing Openat Tracing (Demo)

TOCTOU – sys_enter (Connect)

syscall enter

User Space
Kernel 5.7.0

Execution Flow

```
ptrace_report_syscall(regs, message)
__secure_computing(struct seccomp_data{regs…})
trace_sys_enter(regs, regs->orig_ax)
```

ptrace

seccomp

```
Syscall Table (x86_64)
…
42 sys_connect
43 sys_accept
44 sys_sendto
…
```

```
long __sys_connect((int fd,
        struct sockaddr __user *uservaddr, int addrlen))
{
    …
    struct filename *tmp;
    ret = move_addr_to_kernel
            (uservaddr, addrlen, &address);
    if (!ret)
        ret = __sys_connect_file
            (f.file, &address, addrlen, 0);

    …
}
```
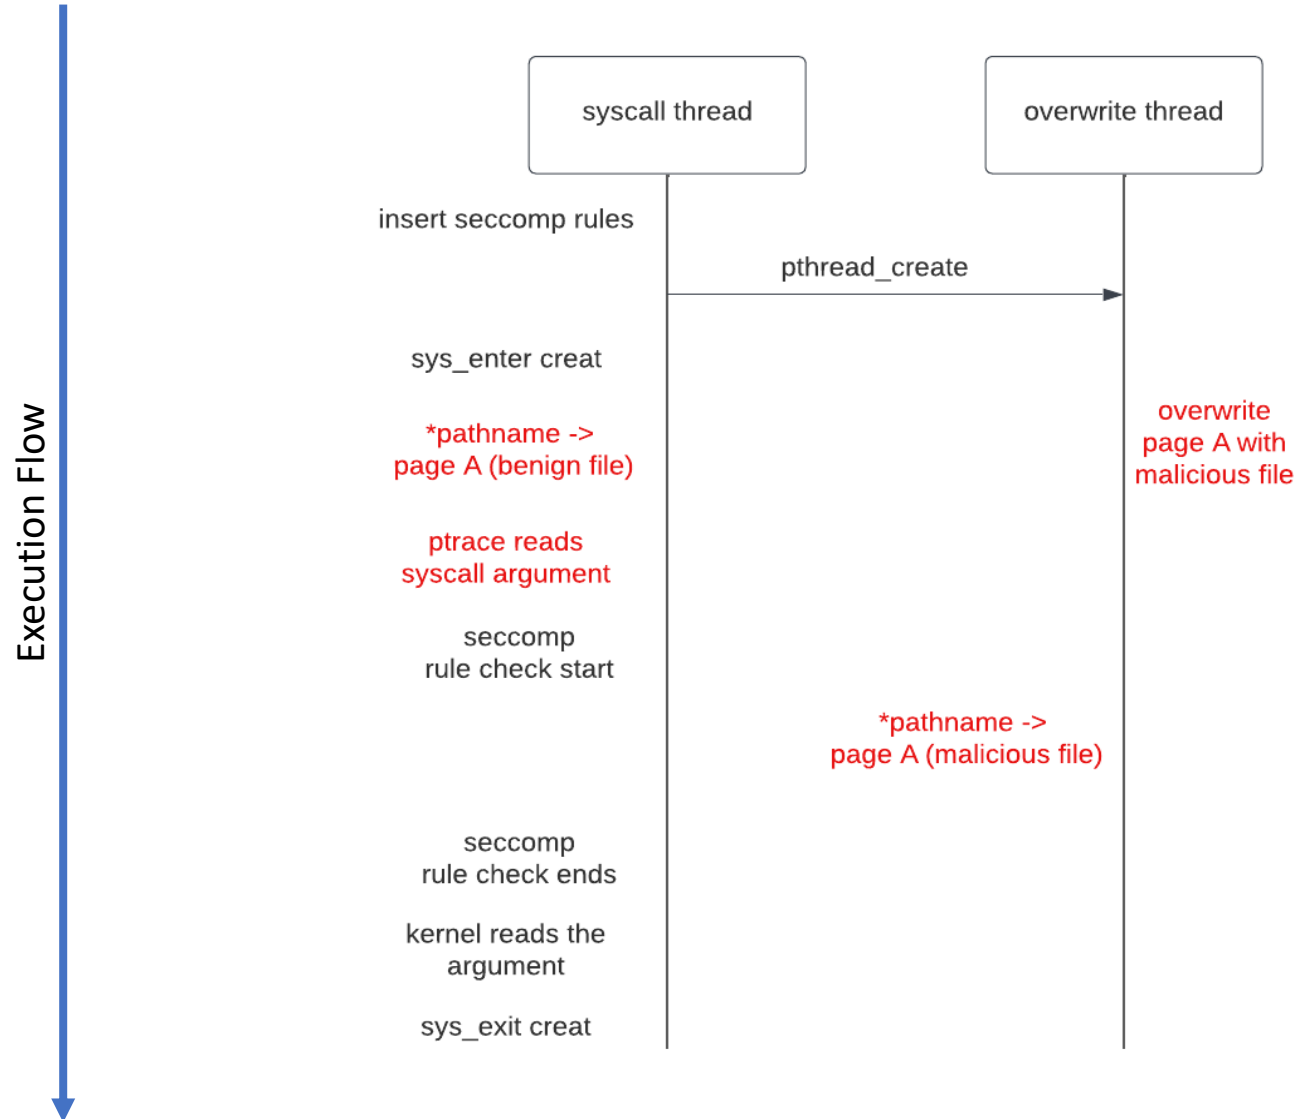
TOU by Linux Kernel

```
trace_sys_exit(regs, regs->ax)
ptrace_report_syscall(regs, message)
```

syscall exit

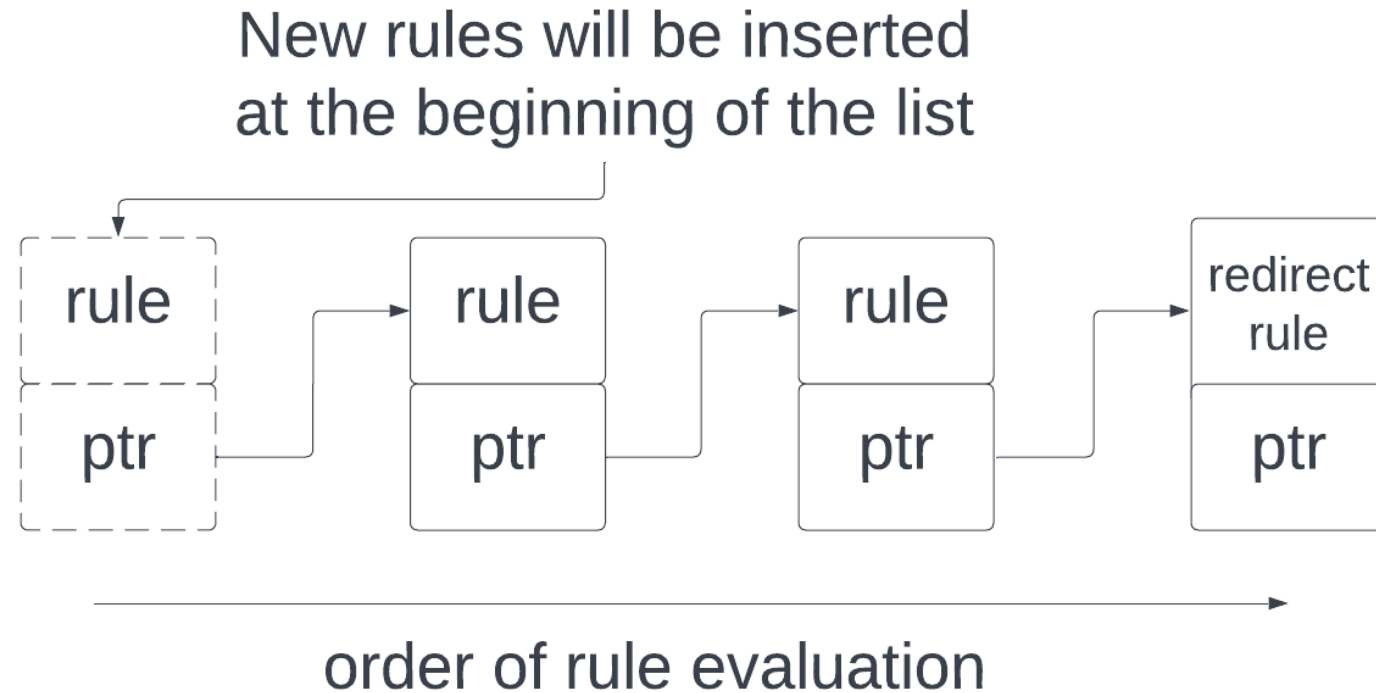**LACEWORK**

Seccomp Introduction

- Kernel level mechanism to restrict syscalls

- Modern sandboxes heavily relies on seccomp

- Developers can write rules to:
  - allow/block certain syscalls
  - allow/block syscalls based on argument values

- These rules can be quite complex ([read more](#))
  - More rules takes more time to compute

- First inserted rules are evaluated last

Attacking Syscall Enter

ptrace + seccomp redirect

- Tracer starts App

New rules will be inserted
at the beginning of the list

| rule | rule | rule | redirect rule |
| ptr | ptr | ptr | ptr |

order of rule evaluation

Exploitation and Mitigations

Tracing location	TOCTOU Exploitation	Mitigations
ptrace @ sys_enter	Seccomp filter insertion	- ptrace + seccomp redirect to start the app. - Inspect seccomp filters already inserted when attaching to a running app
tracepoint @ sys_enter	Unreliable	N/A
tracepoint @ sys_exit	- Blocking syscall (This talk) - Phantom attack v1 (DEFCON 29)	- Compare tracepoint sys_enter and sys_exit args
ptrace @ sys_exit	Same as above	- Deploy all mitigations for ptrace @ sys_enter - Compared the sys_enter and sys_exit syscall args
kprobe @ kernel internal	It depends	Read the kernel copy of the syscall args - LSM (BPF-LSM) - Other interfaces

Key Takeaways

1. Linux kernel tracing can be bypassed reliably
   - Check your security tools

2. Mitigation is complex (workload type and kernel compatibility)
   - Check your security tools' mitigation claims

3. Correlate different data sources

4. Know your normal

- Discussing further?
  - @Xiaofei_REX / rex.guo *NOSPAM* lacework DOT com
  - jzeng04 *NOSPAM* gmail DOT com
  - POC: https://github.com/rexguowork/phantom-attack

LACEWORK

Acknowledgement

- Joel Schopp (Linux kernel / Security)
- Lacework Labs
  - James Condon
  - Greg Foss
  - Chris Hall
  - Jared Stroud
- Falco open source team
  - Leonardo Di Donato
  - Michael Clark
  - Michele Zuccala
  - Luca Guerra
- John Dickson